

# CSE 5194

Implementation of Programming Models and Environments on Modern Systems

Rajarshi Biswas

# Papers:

- [High Performance RDMA-Based MPI Implementation over InfiniBand](#)
  - J. Liu, J. Wu, S. P. Kini, P. Wyckoff, and D. K. Panda
  - Int'l Conference on Supercomputing (ICS '03), June 2003.
- MVAPICH-Aptus: Scalable High-Performance Multi-Transport MPI over InfiniBand
  - M. Koop, T. Jones and D. K. Panda
  - Int'l Parallel and Distributed Processing Symposium (IPDPS), 2008
- Efficient and Truly Passive MPI-3 RMA Using InfiniBand Atomics
  - M. Li, S. Potluri, K. Hamidouche, J. Jose and D. K. Panda
  - EuroMPI 2013, September 2013

# High Performance RDMA-Based MPI Implementation over InfiniBand

J. Liu, J. Wu, S. P. Kini, P. Wyckoff, and D. K. Panda

# Motivation

- MPI has been the de facto standard for writing parallel applications.
- Existing designs of MPI over VIA [and InfiniBand] use send/receive operations for small data messages and control message and RDMA operations for large data messages
- Complexity in hardware implementation and non-transparency to the remote side, send/receive operations do not perform as well as RDMA operations in current InfiniBand platforms.
- This paper we propose a method which brings the benefit of RDMA operations to not only large messages, but also small and control messages.

# InfiniBand Channel and Memory Semantics

- Channel Semantics
  - send/receive operations are used for communication.
    - Sender side, the programmer initiates the send operation by posting a send descriptor
    - Receiver side, the programmer posts a receive descriptor which describes where the message should be put at the receiver side
    - Multiple send and receive descriptors can be posted and they are consumed in FIFO order.
    - Eager Send
- Memory Semantics
  - one-sided and do not incur software overhead at the other side
  - RDMA Read
  - RDMA Write
    - immediate data Variabteion
  - Faster and less overhead

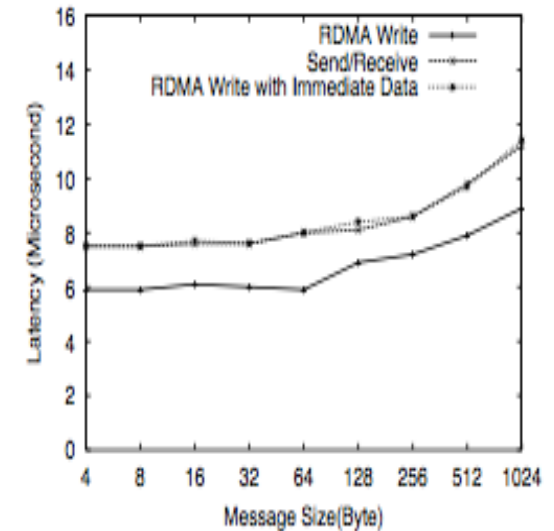


Figure 1: Latency of Send/Receive and RDMA

# Mapping MPI protocols

## Communication Modes

- standard, Synchronous, Buffered, and Ready modes
- Eager protocol (**Send/Receive Based Approach**)
  - Message is pushed to the receiver side regardless of its state.
  - Small message
- Rendezvous
  - A handshake happens between the sender and the receiver via control messages before the data is sent to the receiver side
  - Large message
  - Ideal for zero-copy transfers

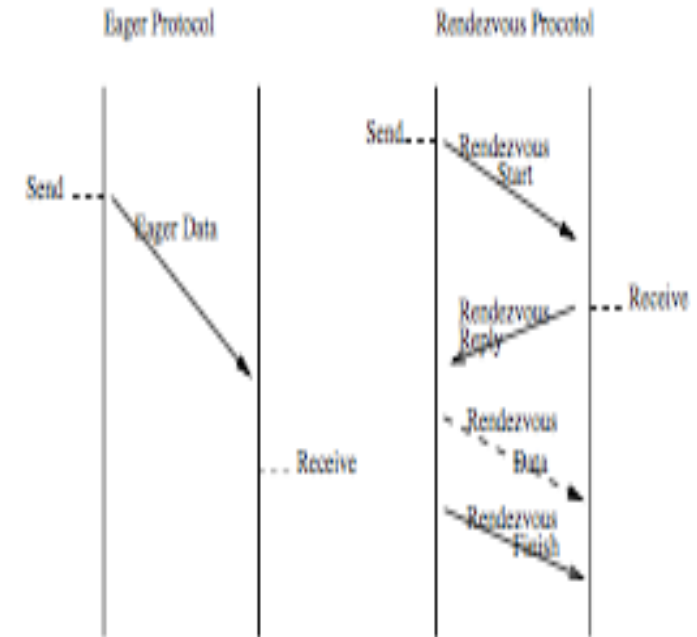


Figure 2: MPI Eager and Rendezvous Protocols

# Mapping MPI protocols

- **Send/Receive Based Approach**

- Eager protocol messages and control messages are transferred using send/receive operations.
- A reliable connection is set up between every two processes
- Buffer pinning and unpinning overhead is avoided by using a pool of pre-pinned,
- A credit based flow control mechanism

- **RDMA-Based Approach**

- RDMA write based approach for Eager protocol and control messages.
- RDMA destination address must be known
- Receiver side must detect the arrival of incoming messages.
- persistent buffer association
  - Reduce the number of RDMA operations
  - persistent correspondence between each buffer at the sender/receiver.
  - reduce the overhead of building RDMA descriptors.
- Receiver always knows exactly where to expect the next message
- Polling time for multiple connection may increase

- **Hybrid Approach**

- RDMA write and send/ receive
- Process only communicates with a subset of all other processes
- RDMA polling set at the receiver side.

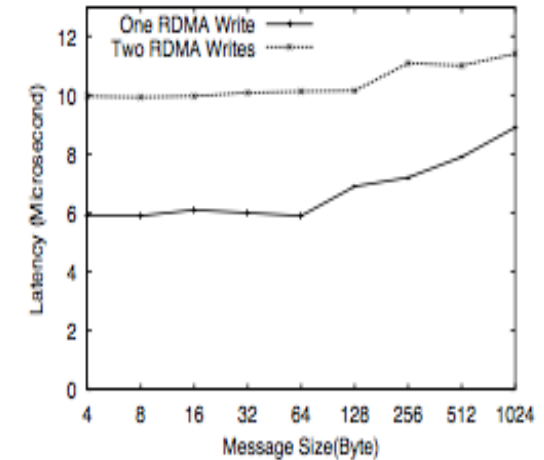


Figure 4: Latency of One RDMA Write versus Two RDMA Writes

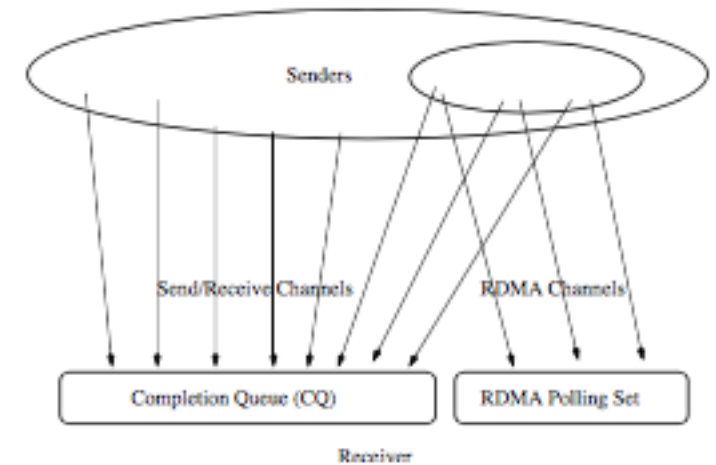


Figure 5: RDMA Polling Set

# Design

- **Basic Structure of an RDMA Channel**

- RDMA channels are unidirectional.
- Set of fixed size, pre-registered buffers at both the sender side and the receiver side.
- Buffers are organized in a Ring
- Sender side, the head pointer is where the next out- going message should be copied
- Tail pointer at the sender side is to record those buffers that are already processed at the receiver side
- Receiver side, the head pointer is where the next incoming message should go
- Tail pointer advances if and only if the current buffer is ready for reuse
- Memory usage




Figure 6: Basic Structure of an RDMA Channel

- **Polling for a Single Connection**

- sender side

- Set data size.
- Check the position of the tail flag. If the value is the same as the primary flag value, use the secondary value. Otherwise, use the primary value.
- Set the head and tail flags.
- Use RDMA write operation to transfer the buffer.

- Receiver polls by performing the following:

- Check to see if the head flag is set. Return if not.
- Read the size and calculate the position of the tail flag. 3. Poll on the tail flag until it is equal to the head flag.
- After processing, the receive side clears the head flag.

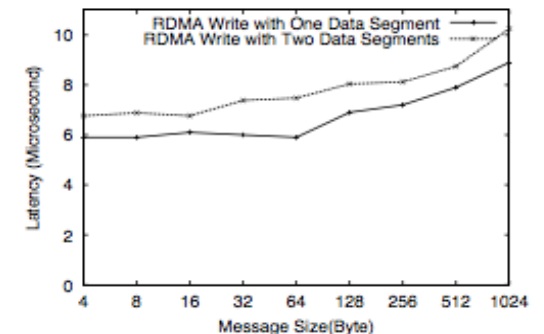


Figure 7: Latency of RDMA Write Gather



# Design

- **Reducing Sender Side Overhead**

- Before buffers can be sent out, descriptors must be allocated and all the fields must be filled
- After the operations are done, completion entries are generated for them in the CQ and the sender side must process them
- Store descriptors with the buffers.
- Unsignalled operations in InfiniBand

- **Flow Control for RDMA Channels**

- **Ensuring Message Order**

- RDMA channel and the send/receive channel.
- Receiver has to poll on both channels to receive messages.
  - Packet Sequence Number Expected Sequence Number

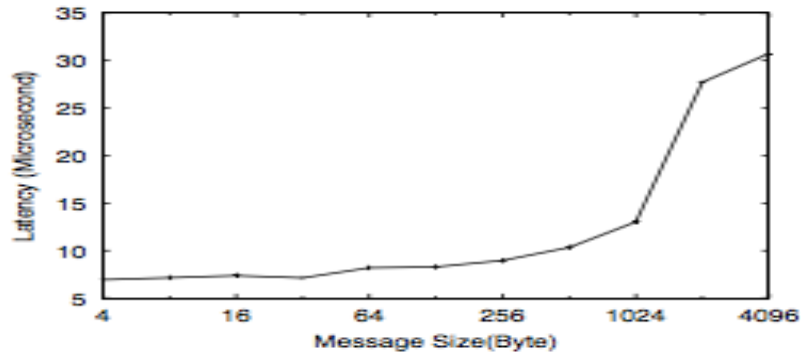
- **Polling Set Management**

- Put first N (N is the size of the RDMA polling set) channels with incoming messages into the RDMA polling set.
- Dynamically manage the polling set for large number of connections


# Performance Evaluation

- **Experimental setup**
- Cluster system
  - consisting of 8 SuperMicro SUPER P4DL6 nodes.
  - Each node has dual Intel Xeon 2.40 GHz processors with a 512K L2 cache at a 400 MHz front side bus.
  - Connected by Mellanox InfiniHost MT23108 DualPort 4X HCA adapter through an InfiniScale MT43132 Eight 4x Port InfiniBand Switch.
  - The HCA adapters work under the PCI-X

# Performance Evaluation



**Figure 9: MPI Latency**

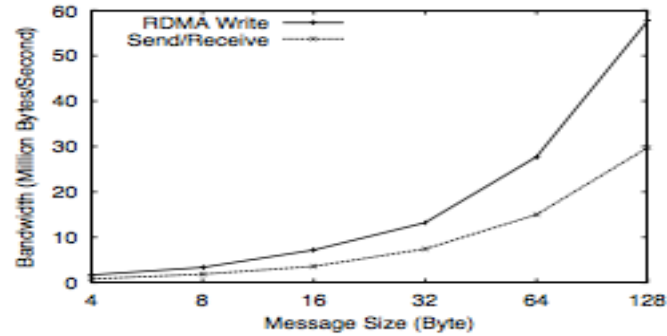


**Figure 10: MPI Bandwidth**

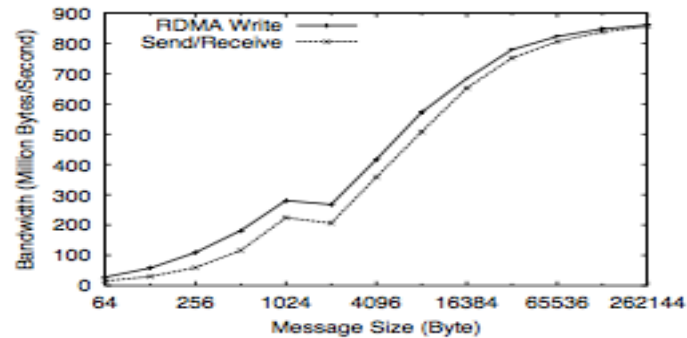
	Latency (us)	Bandwidth (MB/s)
This implementation	6.8	871
Quadrics	4.7	305
Myrinet/GM	7.3	242

- Achieved a 6.8 microseconds latency for small messages
- The peak band- width is around 871 Million Bytes (831 Mega Bytes)/second

# Performance Evaluation




**Figure 12: MPI Bandwidth Comparison (Small Messages)**




**Figure 13: MPI Bandwidth Comparison**


- Improvements of our RDMA- based design by comparing it with the send/receive based design
- *Collective Communication*



**Figure 16: MPI Allreduce Latency (8 Nodes)**




**Figure 17: MPI Broadcast Latency (8 Nodes)**



**Figure 15: MPI Barrier Latency**

# Performance Evaluation



Results for IS, MG, LU, CG, FT, SP and BT programs from the NAS Parallel Bench- mark Suite on 4 and 8 nodes.

Figure 18: NAS Results on 4 Nodes (Class A)




Figure 19: NAS Results on 8 Nodes (Class B)

# Conclusion

- Brings the benefit of RDMA to not only large messages, but also small and control messages
- Better scalability by exploiting application communication pattern and combining send/receive operations with RDMA operations

MVAPICH-Aptus

Scalable High-Performance Multi-Transport MPI over InfiniBand

M. Koop, T. Jones and D. K. Panda

# Problem Statement

- This work seeks to address two main questions:
  - What are the different protocols developed for MPI over InfiniBand and how do they perform at scale?
  - Given this knowledge, can the MPI library be designed to dynamically select protocols to optimize for performance and scalability?



# Message Channels

- Message passing is generally implemented with two modes:
  - Eager Protocol: Small messages (<8K)
  - Rendezvous Protocol: Large messages
- Multiple designs of both protocols have been implemented for InfiniBand
  - Describe and evaluate each of them to determine performance and scalability characteristics

# Eager Channels

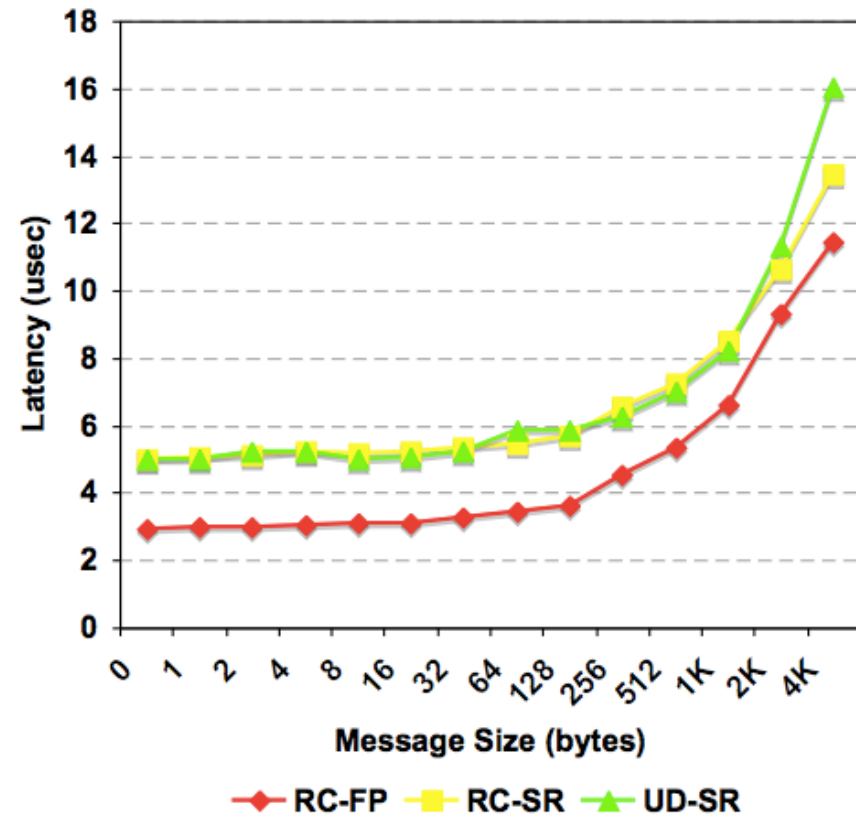
- Reliable Connection Send/Receive (RC-SR)
  - Channel built directly on the channel semantics of the RC transport of InfiniBand
  - Use of the Shared Receive Queue (SRQ) allows pooling of receive buffers to achieve better scalability
- Reliable Connection Fast Path (RC-FP)
  - Current adapters only reach their lowest latency using RDMA Write operations
  - This approach uses paired queues and last-byte polling to achieve low latency (at the cost of memory usage)
- Unreliable Datagram Send/Receive (UD-SR)
  - Built on the channel semantics of the UD transport of InfiniBand
  - Must take care of reliability, however, it is very scalable

# Rendezvous Channels




- Reliable Connection RDMA (RC-RDMA)
  - Using this method an RDMA write operation is used to write directly into the application buffer without intermediate copy operations
- Unreliable Datagram Zero-Copy (UD-ZCopy)
  - Using a pool of QPs and a novel approach, data can be transferred over UD -- preventing the requirement that RC connections be created
- Copy-Based Send
  - Negotiate buffer availability, but then use the eager channels to push the data to the receiver

# Performance: Eager Latency

- Classic ping-pong latency test (osu\_latency)
- *RC-FP delivers lowest latency*
- RC-SR and UD-SR perform similarly until 2K and beyond where UD-SR requires software packetization



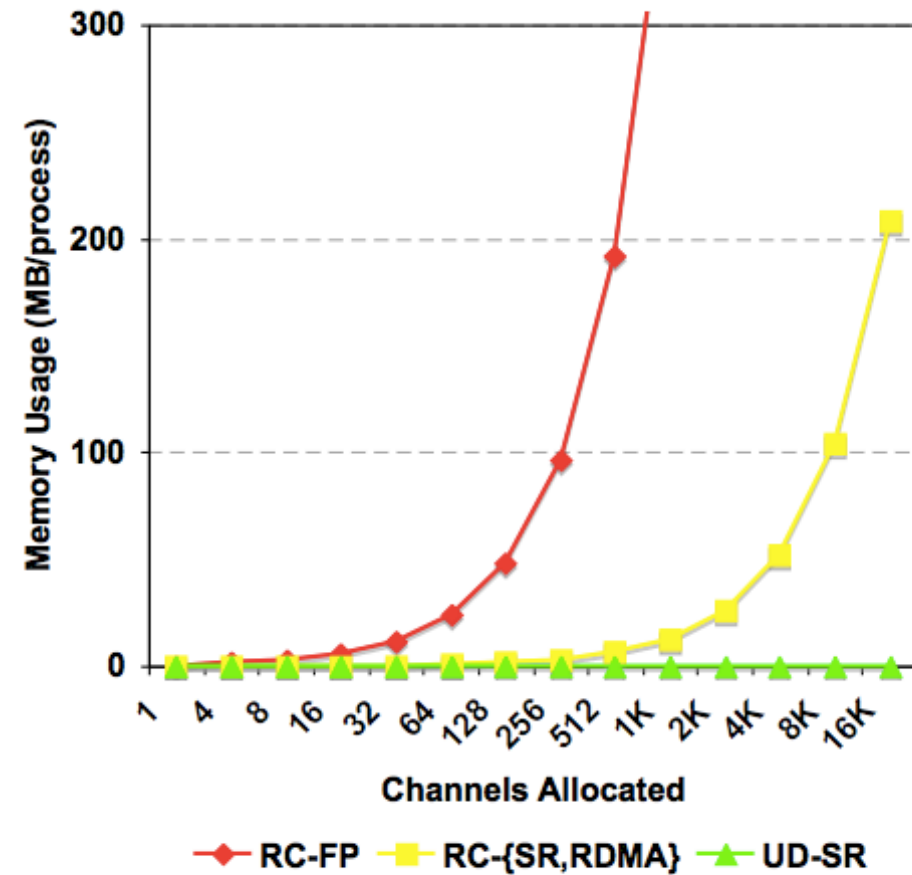
# Performance: Bandwidth





- Throughput for RC-based channels performs poorly when the number of communicating pairs increases
- UD-SR remains scalable in performance

# Scalability: Memory Usage

- *RC-FP* requires a significant amount of memory resources
- *RC-SR* is much more scalable in memory, but can still have issues at scale
- UD-SR remains very scalable with near-constant memory usage



# Scalability: Latency



- Due to the memory polling used in RC-FP only a few channels can be allocated before latency increases
- The InfiniBand HCA has only a limited number of QPs that can be active in the on-card cache

# Summary

Type	Channel	Transport	Latency	Throughput	Scalability
Eager	RC-SR	RC	Good	Fair	Fair
	RC-FP	RC	Best*	Good	Poor
	UD-SR	UD	<2K, Good >2K, Poor	< 2K, Best > 2K, Poor	Best
Rendezvous	RC-RDMA	RC	-	Best	Fair
	UD-ZCopy	UD	-	Good	Best
	Copy-Based	RC/UD	-	Poor	-

No eager or rendezvous channel has all of the desired features



# MVAPICH-Aptus Design

- As seen from the previous evaluation results, no single channel for either eager or rendezvous is always best
- General Goal:
  - Use a combination of message channels and transports to optimize for **performance** and **scalability**
- Design Challenges:
  - *When should a channel be **created**? ■ When should a channel be **used**?*

# Channel Allocation

- Some channels perform well when only a limited number of them are created, but quickly deteriorate
  - RC Transports (RC-SR/RC-FP/RC-RDMA)
    - Each RC connection requires additional memory usage
    - Cache on HCA can be overflowed quickly
  - RC-FP
    - Too many channels increases polling time
    - Memory scalability is poor
- Strategy:
  - Create up to a configurable number of channels of each type ■
    - 16 RC QPs
    - 8 RC-FP connections
  - Setup after a certain number of “qualified” messages are transferred

# Channel Usage

- As found earlier, some channels also perform differently given message size and other features
- We allow a flexible form of matching when sending a message:
- Take the first match where both the conditional is true and the channel is allocated to the destination peer

**Sample  
Configuration**

```
MSG_SIZE <= 2048, RC-FP,  
MSG_SIZE <= 2008, UD-SR,  
MSG_SIZE <= 8192, RC-SR,  
MSG_SIZE <= 8192, UD-SR,  
TRUE, RC-RDMA,  
TRUE, UD-ZCopy,  
TRUE, Copy-Based
```

# Evaluation


- MVAPICH
- Evaluated Configurations:

	RC-SR	RC-RDMA	RC-FP	UD-SR	UD-ZCopy
<b>RC</b> MVAPICH0.9.9	Available	Available	Available		
<b>UD</b> MVAPICH-UD				Available	Available
<b>Aptus</b>	Available	Available	Available	Available	Available

## Experimental Testbed:

- 70 node, 560-core InfiniBand Linux cluster
- Dual 2.3GHz “Clovertown” quad-core processors
- Mellanox MT25208 DDR HCA
- OpenFabrics OFED 1.2
- We evaluate the following application benchmarks
  - NAS Parallel Benchmarks: CFD application kernels
  - NAMD: Molecular dynamics application
  - SMG2000: Multigrid solver (ASC Benchmark)
- In addition to collecting the wallclock performance measurement, we also evaluate other characteristics:
  - Channels created
  - Message and data volume over each channel

# Performance Results

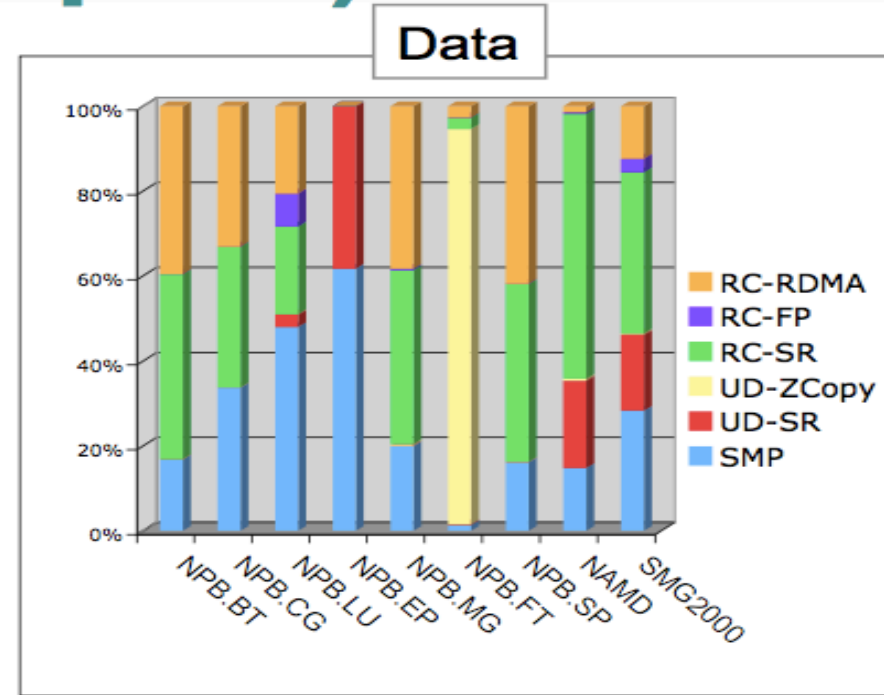



	SMP	UD-{SR,Zcopy}	RC-{SR,RDMA}	RC-FP
NPB.BT	4.11	20.17	10.60	7.88
NPB.CG	3.00	6.94	2.94	2.94
NPB.EP	3.00	6.00	0.00	0.00
NPB.FT	7.00	504.00	16.00	8.00
NPB.MG	4.31	9.00	5.63	5.63
NPB.LU	3.75	7.06	2.23	2.23
NPB.SP	4.11	20.17	10.62	7.88
NAMD	6.30	120.80	16.47	8.00
SMG2000	4.25	120.19	16.34	8.00

- In all results we see that the hybrid UD/RC design is able to outperform or match either mode used exclusively
  - 512/484 processes

Breakdown shows Aptus dynamically has setup the fewest channels needed

# Channel Volume (Aptus)



Breakdown of message transfers by channel show *good utilization* of “expensive” channels, despite allocating only a few of them

# Conclusions

- As clusters continue to scale, the MPI library must be scalable in memory as well as performance
- Previously a UD-based MPI showed superior scalability, but lower performance in some applications
- In this work we bridge the gap between RC and UD designs
- We are working towards
  - Looking into the new eXtended Reliable Connection (XRC) transport provided in ConnectX adapters
  - Release of the Aptus (UD/RC) design in an upcoming version of MVAPICH
  - Investigate support for dynamic communication patterns

# Efficient and Truly Passive MPI-3 RMA Using InfiniBand Atomics

M. Li, S. Potluri, K. Hamidouche, J. Jose and D. K. Panda



# Motivation

- **MPI Remote Memory Access (RMA) Model**
- Minimizing communication overheads is key as applications scale to millions of processes/cores
- RMA model offers an alternative to Send/Recv based message passing model
  - Communication Epochs
    - Period between 2 synchronizations
    - One-sided communication
    - Windows area
- Promises better latency hiding, asynchronous progress and reduced synchronization overheads
- MPI-3 offers several extensions to provide more flexibility

## MPI-3 RMA Passive Synchronization

- RMA offers flexible synchronization alternatives
  - Active: Fence and Post-Wait/Start-Complete
  - Passive: Lock/Unlock, Lock\_all/Unlock\_all
  - Shared/Exclusive (Lock/Unlock) and (Only Shared) (Lock\_all/Unlock\_all)
- Passive synchronization does not require involvement of target process
  - Less synchronization
  - Better overlap
- However, current implementations are based on two-sided operations
- Desirable to have a truly one-sided design offering
  - Performance (no remote polling)
  - Fairness (FIFO)

# Problem Statement

- Can a truly passive locking mechanism be designed for InfiniBand Clusters ?
- How can this design provide :
  - Performance (no remote Polling) - Fairness (FIFO => no starvation)
- Can the new locking mechanism benefits the performance of applications ?

# Existing Passive Synchronization Semantics over IB

**IB**

	Shared	Exclusive	Limitations
State-of-the art MPI Libraries	Send/Recv	Send/Recv	Restrict asynchronous progress
Jiang et.al (Compare_and_swap)	Atomics	Atomics	High network Traffic due to remote polling
Jiang et.al (MCS based)	--	Atomics/Put	Shared mode of locking is not handled
Santhanaraman et.al	Send/Recv	Atomics	Restrict asynchronous progress. High network Traffic

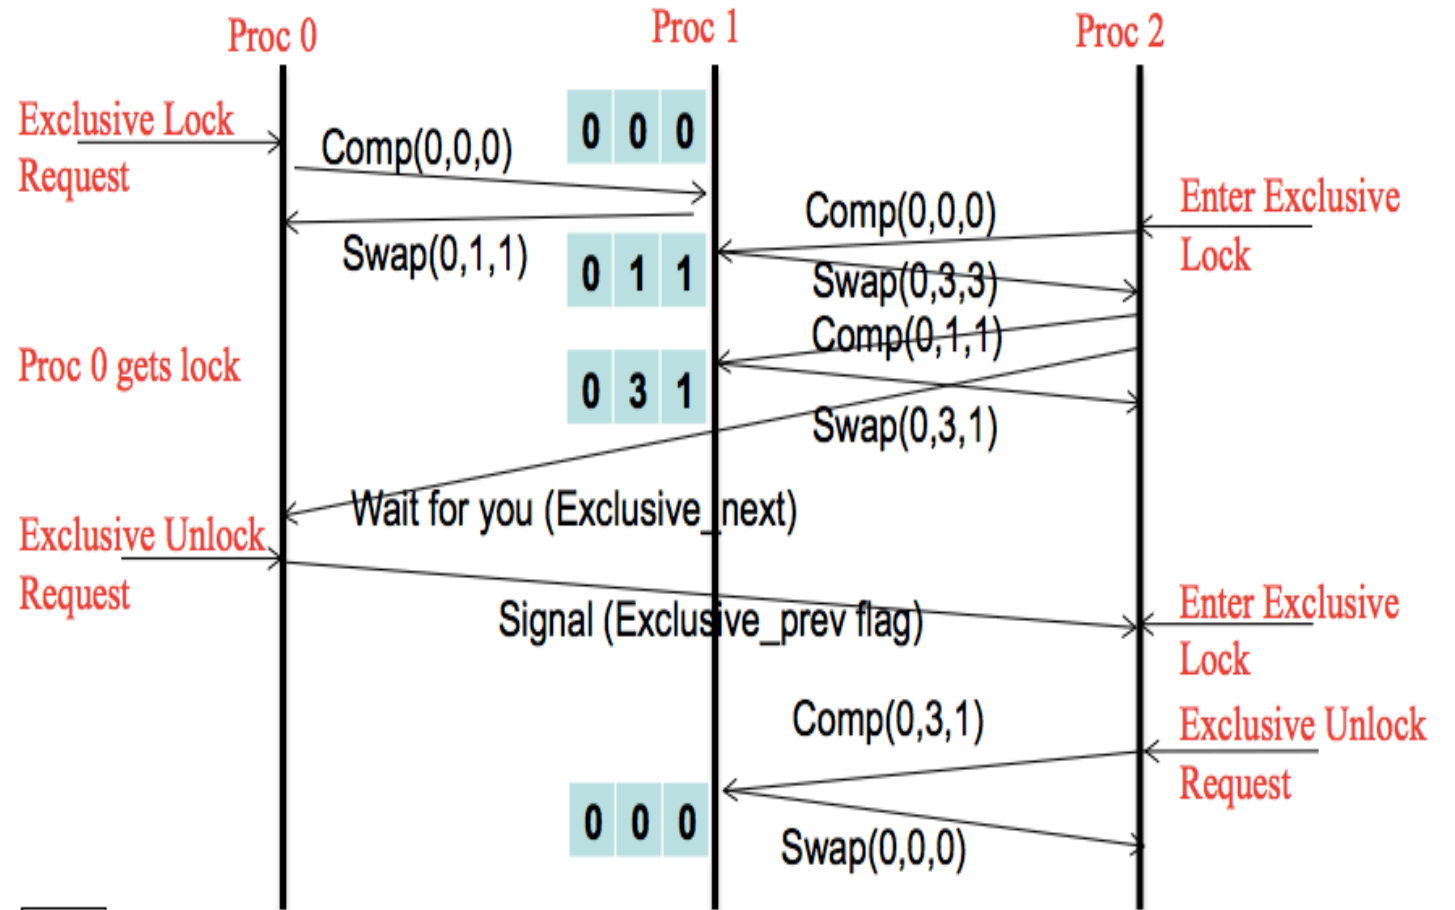
# Efficient and Truly Passive Synchronization scheme

- Lock Data Structures
- Our locking mechanism depends on IB atomics to implement shared and exclusive mode of locking
- IB requires 64 bits buffer for atomic operations
- This 64 bits region is divided into three parts to handle different lock requests
  - Shared Counter: count of the processes that own or have requested a shared lock
  - Exclusive Tail: rank of the process which is tail of the distributed queue
  - Exclusive Head: rank of the process which is head of the distributed queue
- In order to handle all possible lock requests, a distributed lock queue is maintained to ensure FIFO and avoid remote polling
- Data structures to implement the distributed lock queue:
  - – Wait-for array: used when shared lock comes after exclusive lock. This exclusive lock knows the list of processes that request shared lock after it
  - – Signal-to array: used when shared lock comes after exclusive lock. This exclusive lock wakes up pending processes that are waiting for the shared lock
  - – Exclusive-next: two element integer array. Used by processes requesting exclusive lock to form a distributed lock queue
  - – Exclusive-prev: one integer flag. Used by a process unlocking an exclusive lock to wake up another process waiting for an exclusive lock



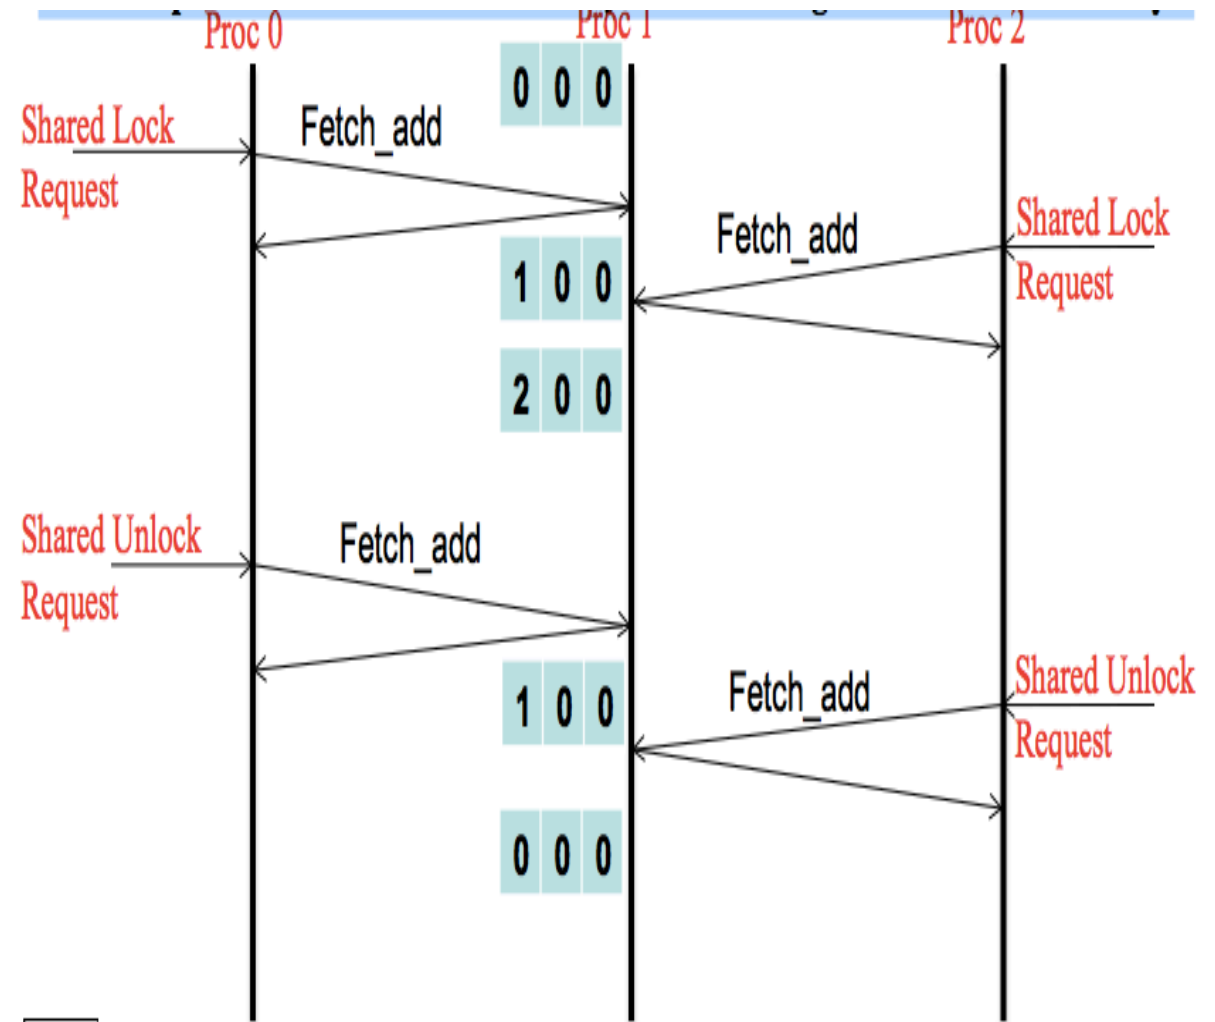
# Exclusive Locking Only

- RDMA operations: `compare_and_swap` and `Put`
- Lock requests are ordered in distributed queue
- Exclusive locks are granted in FIFO order



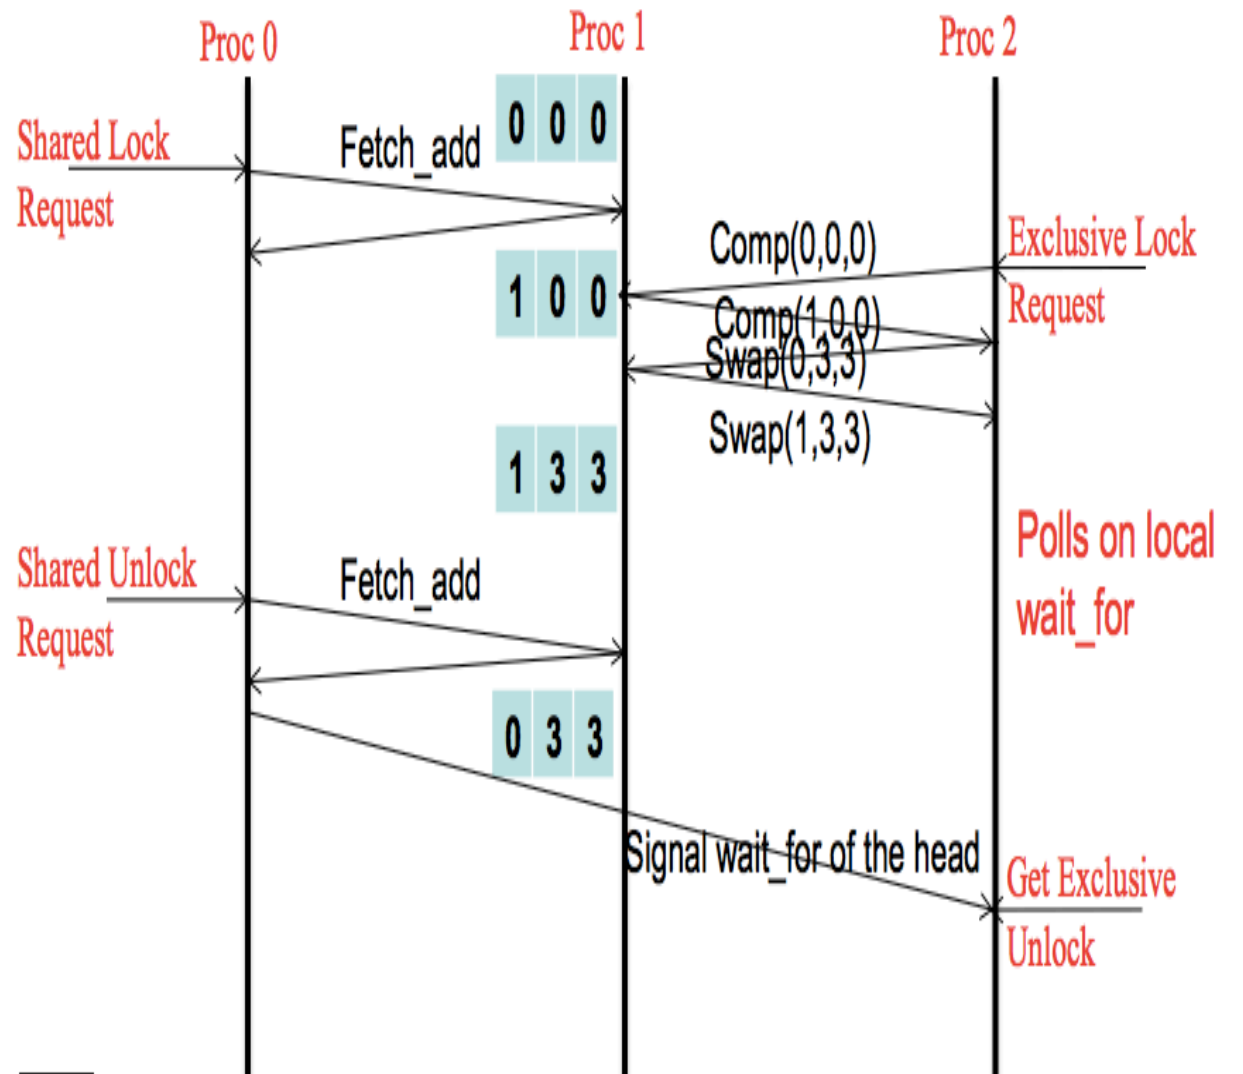
# Shared Locking Only

- Atomic operation : Fetch\_and\_add. To decrement we add the MAX value
- Each process requires shared lock is able to get it after its atomic operation completes
- Each process releases shared lock by decrementing shared lock counter by 1




# Interleaved Shared and Exclusive Locking

- Shared followed by exclusive lock: Process gets exclusive lock after all previously granted shared locks have been released.
- Exclusive followed by shared lock: Process gets shared lock after the previous exclusive lock releases its lock




# Intra-node Locking Design

- For intra-node locking, native loopback that needs a number of queue pairs  $(p + (p * (p - 1)) / 2)$  is not an efficient implementation ( $P$  = number of process on a node)
- If the lock-unlock 64 bits data structures are allocated in the shared memory region, the number of queue pairs used is decreased from  $(p + (p * (p - 1)) / 2)$  to  $p$ 
  - Based on the intra-node locking design, if one process wants to acquire a lock from other process in the same node, it issue atomic operation to itself (loopback)
  - The locking/unlocking mechanisms are the same as discussed earlier



$(P + (P * (P - 1)) / 2)$  QPs



$P$  QPs



# Lock\_All/Unlock\_All Implementation

- **Lock\_all and Unlock\_all introduced in MPI-3 use only shared lock.**
- **In our design, they are implemented based on the lock/unlock mechanism**
- **discussed eariler.**
- **If MPI\_MODE\_NOCHECK is used, then they are implemented as No\_Op**
- **Inside Lock\_all function, call win\_lock is explicitly called for every processes in the communicator**
- **For Unlock\_all, the same mechanism is used to call unlock for every process in the communicator**

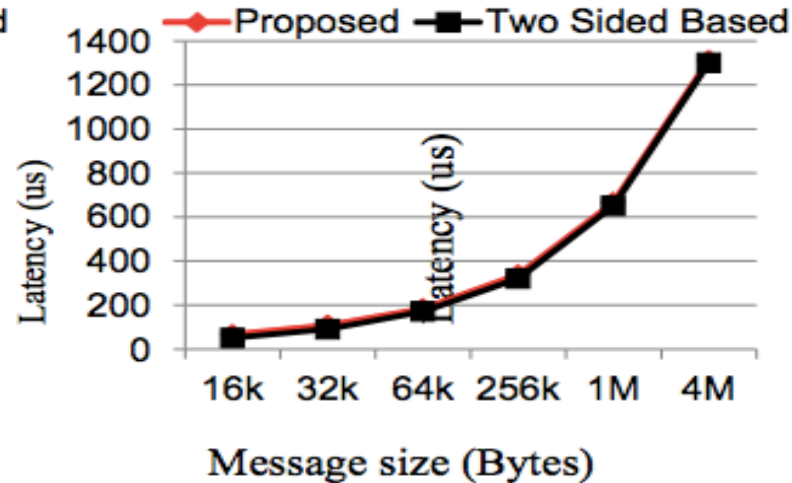
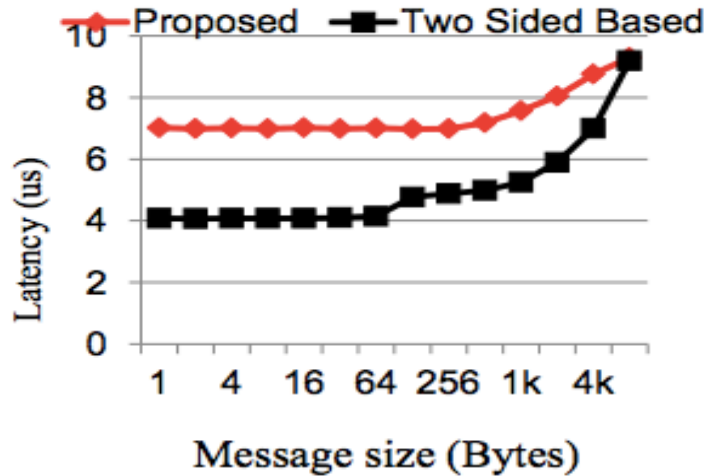
# Performance Evaluation

## Experimental Setup


- **Cluster A**
  - – **Xeon Dual quad-core processor (2.67 GHz) with 12GB RAM**
  - – **Mellanox QDR ConnectX HCAs (32 Gbps data rate) with PCI\_Ex Gen2 interface**
- **• Software stack**
  - – **Implemented on MVAPICH2-1.9 will be in future releases**
  - <http://mvapich.cse.ohio-state.edu> Latest releases : **MVAPICH2-2.0a**
  - High Performance open-source MPI Library for InfiniBand, 10Gig/iWARP, and RDMA over Converged Enhanced Ethernet (RoCE)
    - – MVAPICH (MPI-1) ,MVAPICH2 (MPI-2.2 and MPI-3.0), Available since 2002
    - – MVAPICH2-X (MPI + PGAS), Available since 2012
    - – Used by more than 2,077 organizations (HPC Centers, Industry and Universities) in 70 countries

# MPI\_Get with Lock-Unlock

One MPI\_Get Latency


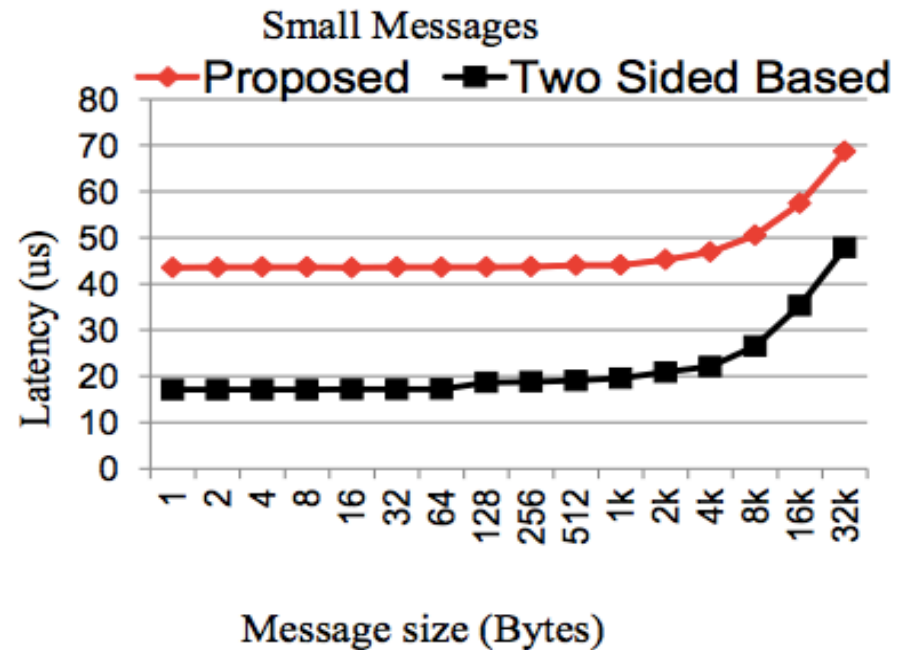


Eight MPI\_Get Latency




- For one MPI\_Get latency:
  - Small messages: atomic based design incurs an overhead compared to two-sided based design : two-sided design coalesces the 3 operations in one message
  - Large messages: Amortized the overhead and have similar performance
- For eight MPI\_Get latency, the overhead is amortized and we see similar performance with both designs

# MPI\_Get with Lock\_all-Unlock\_all




- Same trend for small messages
- The design could benefit large messages by asynchronously issuing lock/unlock requests from different processes

# Overlap Benchmark




Communication Overlap-Lock



Communication Overlap-Lock\_all

- Achieves **almost optimal** computation/communication **overlapping**

# Splash LU Kernel



- This modified version of Splash LU Kernel does dense LU factorization
- Our design outperforms the two-sided approach by a factor of **49%** and **35%** on 4 and 32 processes

# Conclusion

- **Proposed Locking mechanism to implement both shared and exclusive lock with RDMA InfiniBand Atomics:**
- **- No remote polling - FIFO order.**
- **Show optimal computation communication overlap**
- **Demonstrated up to 49% improvement using Splash LU Kernel**
- **Evaluate our designs with more applications/systems**
- **Provide RDMA based-designs for MPI-3 RMA over IB**

# References:

- <http://nowlab.cse.ohio-state.edu/static/media/publications/abstract/liuj-ics03.pdf>
- <http://nowlab.cse.ohio-state.edu/static/media/publications/slide/koop-ipdps08.pdf>
- [http://web.cse.ohio-state.edu/~panda.2/5194/slides/5c\\_eurompi13\\_rma.pdf](http://web.cse.ohio-state.edu/~panda.2/5194/slides/5c_eurompi13_rma.pdf)



# Thank You



- Network-Based Computing Laboratory
  - <http://nowlab.cse.ohio-state.edu/>
- MVAPICH Web Page
  - <http://mvapich.cse.ohio-state.edu/>